

Amit Gohel

amitgohel2002@gmail.com | +91 70166 79944 | amitgohel.dev | linkedin.com/in/amitgohel-ai | github.com/amit-gohel

Professional Summary

AI/ML Engineer with 2+ years delivering Python-based enterprise AI in production. Built RAG chatbots at 95% MRR handling 50,000+ queries/month, agentic LangChain pipelines on multi-tenant platforms, and on-premise document processing for sensitive data. Hands-on with multi-model LLM orchestration across GPT-4o, Anthropic Claude, and Gemini 2.5 Pro, retrieval stack optimization on Milvus (111% MRR improvement from FAISS), and production-scale embedding evaluation (6+ models, BGE-M3 + reranker). Microsoft Azure AI Engineer Associate certified.

Work Experience

AI Engineer, Vivansh Infotech *Ahmedabad, India* Aug 2025 – Present

- Upgraded permission-aware RAG chatbot to fully **agentic architecture using LangChain/LangGraph** for multi-tenant EHS platform serving **500+ concurrent users, 1,000+ queries/day**; integrated Milvus, PostgreSQL, and 10+ platform APIs for autonomous data retrieval and workflow triggers
- Enforced **250+ granular permissions per query** via 2-stage retrieval: dedicated Milvus permission index cross-referenced with PostgreSQL, with GPT-4o handling access control, query rewriting, and conversation resolution across 10,000+ chunks; **80% MRR**
- Built **AI-powered form builder** converting PDF forms to digital replicas in ~20 seconds (**200+ forms in month 1**); Pydantic-validated intermediate JSON schema enabled reliable LLM-to-platform conversion across multi-model pipeline (GPT-4o, Qwen3-Coder-480B, GPT-4o-mini, Anthropic Claude) selected via empirical benchmarking

AI/ML Developer, Silver Touch Technologies Ltd *Ahmedabad, India* Jun 2024 – Jul 2025

- Rebuilt production RAG chatbot infrastructure serving **50,000+ queries/month** across 10+ chatbot instances; migrated from in-code Llama-2-7b to VLLM-served API and FAISS to Milvus, improving **MRR by 111% (45% to 95%)**
- Led vector DB evaluation (Milvus vs. Qdrant vs. FAISS) on production-scale data; tested 6+ embedding models (BGE-M3, BGE-small/base/large, Stella, SPLADE) and built custom retrieval with BGE-M3 + reranker
- Reduced response latency by **50% (16s to 8s)** through query preprocessing with intent identification and retrieval pipeline optimization
- Built **on-premise OCR and document pipeline** for sensitive government documents (200+ page PDFs); merged DocTR + DocLink via Llama-3.1-Nemotron-70b, achieving **~90% entity extraction accuracy** on Aadhaar, PAN, and GST numbers

Research Intern, Indian Institute of Technology (IIT) Delhi *New Delhi, India* Jan 2024 – May 2024

- Collected 12,400 labeled samples from 20 participants; evaluated **1,000+ model configurations** across 5 architecture families (Dense, CNN, LSTM, BiLSTM, hybrids) using TensorFlow/Keras – data-driven model selection at scale; collaborated with IISc Bangalore and Howard University researchers

Technical Skills

- Languages and Frameworks:** Python, C++, SQL; LangChain, LangGraph, PyTorch, TensorFlow/Keras, Hugging Face Transformers, FastAPI, Pydantic
- Vector DBs and Retrieval:** Milvus, Qdrant, FAISS; BGE-M3, Cohere Embed/Rerank, SPLADE, OpenAI Embeddings
- LLMs and Techniques:** GPT-4o, Anthropic Claude, Gemini 2.5 Pro, Llama, Qwen3-Coder; RAG (dense/sparse retrieval), Agentic AI, Prompt Engineering (COT, ReAct), NLP, Fine-tuning (PEFT), CNN, LSTM, BiLSTM
- Infrastructure and Tools:** AWS (EC2, Lambda, S3, Bedrock, SAM), Azure, GCP, Docker, Git, PostgreSQL, VLLM, OpenAI/Anthropic/Gemini APIs, Together AI, Selenium, FFmpeg

Education

B.E. Computer Engineering, Vishwakarma Government Engineering College, Ahmedabad Jun 2020 – Jun 2024

Certifications and Achievements

- Microsoft Azure AI Engineer Associate** – Certified in designing and implementing Azure AI solutions
- Samsung Solve for Tomorrow** – Top 10 among 70,000+ participants nationwide; featured on CNN News18
- Smart India Hackathon 2023** – Top 5 among 44,000+ teams nationwide (MIT Karnataka)
- Competitive Programming:** 370+ problems solved in Python