

Amit Gohel

amitgohel2002@gmail.com | +91 70166 79944 | amitgohel.dev

linkedin.com/in/amitgohel-ai | github.com/amit-gohel

Professional Summary

AI Engineer with 2+ years shipping production agentic systems and RAG infrastructure. Built multi-agent LangGraph + LangChain pipelines serving 50,000+ queries/month with 250+ real-time permission checks; led FAISS→Milvus migration (MRR 45%→95%); orchestrated multi-model workflows (GPT-4o, Gemini, Qwen3) in production. Python, FastAPI. Microsoft Azure AI Engineer Associate certified.

Work Experience

AI Engineer, Vivansh Infotech *Ahmedabad, India* Aug 2025 – Present

- Upgraded permission-aware RAG chatbot to a fully agentic architecture using **LangGraph** (stateful multi-agent graph) + **LangChain** (tool-calling over 10+ APIs) on a multi-tenant EHS platform; **Milvus** + **PostgreSQL** backend serving **500+ concurrent users**, autonomous live data retrieval and workflow execution
- Enforced **250+ permissions per query** via 2-stage Milvus retrieval (permission index + PostgreSQL cross-check); GPT-4o for access control and query rewriting across 10,000+ chunks; **80% MRR**
- Built **AI form builder** converting PDFs to digital replicas in ~20 seconds (**200+ forms in month 1**); multi-model pipeline (GPT-4o, Qwen3-Coder-480B, GPT-4o-mini) with Pydantic-validated JSON schema
- AI video course pipeline** (1 week → 10 min, 99% reduction): Milvus + Cohere reranker, Gemini 2.5 Pro, ElevenLabs narration, serverless FFmpeg on AWS Lambda

AI/ML Developer, Silver Touch Technologies Ltd *Ahmedabad, India* Jun 2024 – Jul 2025

- Rebuilt production RAG platform serving **50,000+ queries/month** across 10+ instances – migrated Llama-2-7b to VLLM + FAISS to Milvus; MRR improved **111% (45%→95%)**
- Led evaluation of Milvus vs. Qdrant vs. FAISS + 6 embedding models (BGE-M3, Stella, SPLADE and variants) on production data; built BGE-M3 + reranker retrieval pipeline
- Reduced response latency **50% (16s→8s)** via query preprocessing with intent classification and retrieval optimization
- Built **on-premise OCR pipeline** for sensitive government documents (200+ page PDFs) using DocTR + DocLink merged via Llama-3.1-Nemotron-70b; ~90% entity extraction on Aadhaar, PAN, GST

Technical Skills

- Programming Languages:** Python, C++, SQL
- AI/ML Frameworks:** LangGraph, LangChain, PyTorch, TensorFlow/Keras, Hugging Face Transformers, LLMs (GPT-4o, Gemini 2.5 Pro, Llama, Qwen3-Coder)
- Agentic AI and RAG:** Multi-agent systems, tool calling, RAG pipelines, prompt engineering, eval design, agentic orchestration, NLP
- Vector Databases and Retrieval:** Milvus, FAISS, Qdrant, OpenAI Embeddings, BGE-M3, Cohere Embed/Rerank, SPLADE
- Infrastructure and DevOps:** FastAPI, VLLM, Docker, PostgreSQL, AWS (EC2, Lambda, S3, SAM), Pydantic, Git
- Tools and APIs:** OpenAI API, Google Gemini API, Anthropic API, Together AI, ElevenLabs, FFmpeg

Education

Bachelor of Engineering – Computer Engineering Jun 2020 – Jun 2024

Vishwakarma Government Engineering College, Ahmedabad, India

Certifications and Achievements

- Microsoft Azure AI Engineer Associate** – Certified in designing and implementing Azure AI solutions
- Samsung Solve for Tomorrow** – Top 10 finalist among 70,000+ participants nationwide; Voicer, wearable ISL-to-speech AI; featured on CNN News18
- Smart India Hackathon 2023** – Top 5 among 44,000+ teams nationwide (Hardware Edition, MIT Karnataka)
- Competitive Programming:** 370+ problems solved in Python