

Amit Gohel

amitgohel2002@gmail.com | +91 70166 79944 | amitgohel.dev | linkedin.com/in/amitgohel-ai | github.com/amit-gohel

Professional Summary

AI/ML Engineer with 2+ years building production ML engineering platforms at scale. Designed and shipped inferencing pipelines, feature engineering systems, and multi-model orchestration architectures serving 50,000+ queries/month; improved model retrieval MRR from 45% to 95% through systematic embedding evaluation and pipeline optimization. Strong Python, AWS, SQL, and cloud ML deployment experience. Microsoft Azure AI Engineer Associate certified.

Work Experience

AI Engineer, Vivansh Infotech *Ahmedabad, India* Aug 2025 – Present

- Upgraded permission-aware RAG chatbot to fully **agentic architecture using LangChain/LangGraph** for multi-tenant EHS platform serving **500+ concurrent users, 1,000+ queries/day**; integrated Milvus, PostgreSQL, and 10+ platform APIs for autonomous data retrieval and workflow triggers
- Enforced **250+ granular permissions per query** via 2-stage retrieval: dedicated Milvus permission index cross-referenced with PostgreSQL, with GPT-4o handling access control, query rewriting, and conversation resolution across 10,000+ chunks; **80% MRR**
- Built **AI-powered form builder** converting PDF forms to digital replicas in ~20 seconds (**200+ forms in month 1**), bulk processing 100+ PDFs in 5–10 minutes; replaced 1+ hour of manual work per form via Pydantic-validated intermediate JSON and multi-model pipeline (GPT-4o, Qwen3-Coder-480B, GPT-4o-mini)
- Created **AI video course generation pipeline** reducing production from ~1 week to ~10 minutes (**99% reduction**) via RAG-based video selection (Milvus + Cohere), Gemini 2.5 Pro for visual-audio alignment, ElevenLabs narration, and serverless FFmpeg rendering on AWS Lambda

AI/ML Developer, Silver Touch Technologies Ltd *Ahmedabad, India* Jun 2024 – Jul 2025

- Rebuilt production RAG chatbot infrastructure serving **50,000+ queries/month** across 10+ chatbot instances; migrated from in-code Llama-2-7b to VLLM-served API and FAISS to Milvus, improving **MRR by 111% (45% to 95%)**
- Led vector DB evaluation (Milvus vs. Qdrant vs. FAISS) on production-scale data; tested 6+ embedding models (BGE-M3, BGE-small/base/large, Stella, SPLADE) and built custom retrieval with BGE-M3 + reranker
- Reduced response latency by **50% (16s to 8s)** through query preprocessing with intent identification and retrieval pipeline optimization
- Built **on-premise OCR and document pipeline** for sensitive government documents (200+ page PDFs); merged DocTR + DocLink via Llama-3.1-Nemotron-70b, achieving **~90% entity extraction accuracy** on Aadhaar, PAN, and GST numbers

Research Intern, Indian Institute of Technology (IIT) Delhi *New Delhi, India* Jan 2024 – May 2024

- Evaluated **1,000+ model configurations** across 5 architecture families (Dense, CNN, LSTM, BiLSTM, hybrids) in TensorFlow/Keras for 62-class gesture classification; selected CNN-BiLSTM via leave-one-subject-out cross-validation; collected 12,400 labeled samples from 20 participants

Technical Skills

- Languages and Frameworks:** Python, SQL, C++; PyTorch, TensorFlow/Keras, Hugging Face Transformers, LangChain, LangGraph, FastAPI, Pydantic
- ML Engineering and MLOps:** Feature engineering, inferencing pipelines, model performance tracking, retrieval pipeline optimization; VLLM (model serving), Docker, Git
- Cloud and Infrastructure:** AWS (EC2, Lambda, S3, IAM), Azure (AI Engineer Associate certified); PostgreSQL, on-premise GPU deployment
- LLMs and AI Techniques:** GPT-4o, Gemini 2.5 Pro, Llama, Qwen3-Coder; RAG, Agentic AI, Prompt Engineering, NLP, CNN, LSTM, BiLSTM, GRU
- Vector DBs and Retrieval:** Milvus, FAISS, Qdrant; OpenAI Embeddings, BGE-M3, Cohere Embed/Rerank, SPLADE, Stella

Projects

On-Premise OCR + Document Entity Extraction Pipeline 2024

- Built hybrid OCR pipeline for sensitive government documents (200+ page PDFs, zero cloud APIs): ran DocTR + DocLink per page and used on-premise Llama-3.1-Nemotron-70b to merge outputs; achieved **~90% entity extraction accuracy** on Aadhaar, PAN, and GST numbers; fully self-hosted on client's 250GB server

SAP Analytics Bot – Natural Language to SQL + Auto-Dashboard 2024

- Production Text-to-SQL system letting business users query live SAP-connected databases in plain English; auto-generates dashboard charts from results; handles multi-step clarification, large result preprocessing, and statistical summaries – no SQL knowledge required

Education

B.E. Computer Engineering, Vishwakarma Government Engineering College, Ahmedabad Jun 2020 – Jun 2024

Certifications and Achievements

- Microsoft Azure AI Engineer Associate** – Certified in designing and implementing Azure AI solutions
- Smart India Hackathon 2023** – Top 5 among 44,000+ teams nationwide (Hardware Edition, MIT Karnataka)
- Competitive Programming:** 370+ problems solved in Python